

CMPS 182, Final Exam, Spring 2019, Shel Finkelstein

ANSWERS

Student Name: _____

Student ID: _____

UCSC Email: _____

Part	Max Points	Points
I	42	
II	24	
III	36	
Total	102	

The first Section (Part I) of the Spring 2019 CMPS 182 Final is multiple choice and is double-sided. Answer all multiple choice questions on your Scantron sheet. You do not have to hand in the first Section of the Exam, but you must hand in the Scantron sheet, with your name, email and student id on that Scantron sheet. Please be sure to use a #2 pencil to mark your choices on this Section of the Final.

This separate second Section (Parts II and III) of the Final is not multiple choice and is single-sided, so that you have extra space to write your answers. If you use that extra space, please be sure to write the number of the problem that you're solving next to your answer. Please write your name, email and student id on this second Section of the Exam, which you must hand in. You may use any writing implement on this Section of the Exam.

At the end of the Final, please be sure to hand in your **Scantron sheet** for the first Section of the Exam and also **this second Section of the Exam**, and show your **UCSC id** when you hand them in. Also hand in your **8.5 x 11 "cheat sheet"**, with your name written in the upper right corner of the sheet.

Part II: (24 points, 6 points each)

Question 22: Here are statements creating two relations R and S:

```
CREATE TABLE S (  
  c INT PRIMARY KEY,  
  d INT  
);
```

```
CREATE TABLE R (  
  a INT PRIMARY KEY,  
  b INT,  
  FOREIGN KEY(b)  
    REFERENCES S(c)  
    ON UPDATE CASCADE  
    ON DELETE SET NULL  
);
```

Relation R(a,b) currently contains the four tuples, (0,4), (1,5), (2,4), and (3,5).

Relation S(c,d) currently contains the four tuples, (2,80), (3,81), (4,82), and (5,83).

Indicate all changes that happen to both R and S when the following SQL statements are executed upon the R and S instances that are shown above. That is, you should **ignore changes made by earlier parts** of this question when you answer later parts.

22a): UPDATE S set c=7 WHERE d=82;

Answer 22a):

Changes made to the original R:

(0,4) is updated to (0,7), and (2,4) is updated to (2,7)

Changes made to the original S:

(4,82) is updated to (7,82)

22b): DELETE FROM R WHERE a =2;

Answer 22b):

Changes made to the original R:

(2,4) is deleted

Changes made to the original S:

No changes

22c): DELETE FROM S WHERE c=5;

Answer 22c):

Changes made to the original R:

(1,5) is updated to (1,NULL), and (3,5) is updated to (3,NULL)

Changes made to the original S:

(5,83) is deleted

Question 23a): What are advantages of having Stored Procedures in a database? Give two different clear advantages. If you give more than two, only first two will be graded.

Answer 23a):

Any two of the following three reasons are good:

- **Performance:** Having a stored procedure means that code can be executed directly in the database, without having to move data between the database and the client, and taking advantage of access paths in the database.
- **Reuse:** Once a Stored Procedure has been written, it can be re-used by the author, or by anyone else who's allowed to use it, without having to write it again.
- **Security:** A user may be granted the right to run a Stored Procedure, even though that user is not allowed to access the data that the Stored Procedure uses.

Question 23b): Since Stored Procedures have the advantages that you listed in part a), why are there also other approaches that combine programming language constructs with SQL?

Answer 23b): Any one of the following reasons would be a good answer.

Stored Procedures can't do everything that (for example) a C or Java program can do. For example, they can't handle user interactions, and they don't have access to all the libraries that C and Java programs can employ.

Also, putting more functionality into the database increases the workload of the database, whereas there can be many instances of clients running C/Java programs.

Question 24: In the dashed lines, write **YES** if the statement is correct, and **NO** if the statement is incorrect.

24a) If the value of the attribute salary1 is NULL, then the truth value for the predicate “salary1 = 1000” will be UNKNOWN.

Answer a) **YES**

24b) If the value of the attribute salary2 is NULL, then the truth value for the predicate “salary2 = NULL” will be UNKNOWN.

Answer b) **YES**

24c) If the value of the attribute salary1 is 1000, and the value of the attribute salary2 is NULL, then the truth value for the predicate “salary1 = 1000 AND salary2 = 1000” will be UNKNOWN.

Answer c) **YES**

24d) A tuple contributes to the result of a query only if its WHERE clause evaluates to TRUE, not if its WHERE evaluates to FALSE or UNKNOWN.

Answer d) **YES**

25e) A CHECK constraint is violated if it evaluates to FALSE or UNKNOWN.

Answer e) **NO**

24f) Assume that our database has a table Customers(cid, cname, level, type, age). If there are no customers whose type is ‘ski’, then the result of the following query will always be the empty set.

```
SELECT c.name
FROM Customers c
WHERE c.age > ALL ( SELECT c2.age
                    FROM Customers c2
                    WHERE c2.type = 'ski' );
```

Answer f) **NO**

Question 25: You have a relation `Company_Info(Emp, Dept, Manager)`, with the following non-trivial Functional Dependencies:

$\text{Emp} \rightarrow \text{Dept}$
 $\text{Emp} \rightarrow \text{Manager}$
 $\text{Dept} \rightarrow \text{Manager}$

25a): Is `Company_Info` in 3NF? Justify your answer clearly.

Answer 25a) No, it's not in 3NF. For `Company_Info` to be in 3NF, every FD would have either a) be Trivial, b) have a superkey on its left-hand side, or c) have a right-hand side that's part of some key.

`Emp` is a superkey because its attribute closure Emp^+ is all the attributes of `Company_Info`. (In fact, it's a key.) But what about the third FD, $\text{Dept} \rightarrow \text{Manager}$?

- a) It's not Trivial.
- b) The attribute closure Dept^+ is just $\{\text{Dept}, \text{Manager}\}$, so the left-hand side is not a superkey.
- c) Moreover, `Manager` is not part of any key. Why? Because no FD has `Emp` on the right-hand side, so any superkey would have to include `Emp`. And $\{\text{Emp}, \text{Manager}\}$ is indeed a superkey, since $\{\text{Emp}, \text{Manager}\}^+ = \{\text{Emp}, \text{Manager}, \text{Dept}\}$. But $\{\text{Emp}, \text{Manager}\}$ is not a key, since you can get rid of `Manager` and still have a superkey, $\{\text{Emp}\}$.

Hence `Company_Info` with those FDs is not in 3NF. (Correct answer doesn't have to discuss the first two FDs, just the third.)

25b): In Design Theory, what is the Update Anomaly? Explain why it is a problem by giving a clear example.

Answer 25b) The Update Anomaly refers to having relation `R` (with Functional Dependencies specified) where the relation is not in 3NF. If you're not careful about updates, there could be tuples in `R` that violate an FD because one tuple is updated, but other tuples would also be updated so that the constraint specified by the FD continues to be obeyed.

Here's one example of this; there are many others, including the relation from part a).

If we have the relation `Employees(eid, name, addr, rank, salary_scale)`, together with the FD: $\text{rank} \rightarrow \text{salary_scale}$, then if there are 10,000 employees who all have rank 4, and the `salary_scale` of one of those employees is changed, then the salary scale of the other 9999 employees who have rank 4 would also have to be updated the same way, or else the database would be inconsistent because it wouldn't obey the FD.

That's a lot of work to do, and a lot of redundancy to manage. Separating out the `Employees` relation into two relations, `Employees2(eid, name, addr, rank)` and `Salary_Table(rank, salary_scale)` would prevent Anomalies.

Part III: (36 points, 9 points each)

Some familiar tables appear below, with Primary Keys underlined. **These tables also appear on the last page of the Final, which you can tear off to help you do questions in Part III of the Final.** You don't have to hand in that last page at the end of the Exam.

Products(productID, productName, manuf, normalPrice, discount)

Customers(customerID, custName, address, joinDate, amountOwed, lastPaidDate, status)

Stores(storeID, storeName, region, address, manager)

Days(dayDate, category)

Sales(productID, customerID, storeID, dayDate, paidPrice, quantity)

Payments(customerID, paidDate, amountPaid, cleared)

Assume that no attributes can be NULL, and that there are no UNIQUE constraints.

Assume Referential Integrity as follows:

- Each productID in Sales appears as a productID in Products.
- Each customerID in Sales appears as a customerID in Customers.
- Each storeID in Sales appears as a storeID in Stores.
- Each dayDate in Sales appears as a dayDate in Days.
- Each customerID in Payments appears as a customerID in Customers.

Write legal SQL queries for Questions 25-28. If you want to create and then use views to answer these questions, that's okay, but views are not required unless the question asks for them.

Don't use DISTINCT in your queries unless it's necessary; 1 point will be deducted if you do. Some points will also be deducted for unnecessarily complex queries, including use of GROUP BY when it's not needed.

Question 26: Find the ID, name and address for each customer whose name begins with 'Prof' and who has no payments that have cleared. (cleared is a Boolean attribute in Payments.) Order your result by name in reverse alphabetical order. Duplicates should not appear in your result.

Answer 26:

```
SELECT c.customerID, c.custname, c.address
FROM Customers c
WHERE c.custname LIKE 'Prof%'
      AND NOT EXISTS ( SELECT *
                        FROM Payments pa
                        WHERE pa.customerID = c.customerID
                        AND pa.cleared )
ORDER BY c.custname DESC;
```

Could also write pa.cleared = TRUE. Different tuple variables could be used.
Don't need DISTINCT, because customerID is a key.

Here are some other correct solutions.

```
SELECT c.customerID, c.custname, c.address
FROM Customers c
WHERE c.custname LIKE 'Prof%'
      AND c.customerID NOT IN ( SELECT pa.customerID
                                FROM Payments pa
                                WHERE pa.cleared )
ORDER BY c.custname DESC;
```

```
SELECT c.customerID, c.custname, c.address
FROM Customers c
WHERE c.custname LIKE 'Prof%'
      AND c.customerID != ANY ( SELECT pa.customerID
                                FROM Payments pa
                                WHERE pa.cleared )
ORDER BY c.custname DESC;
```

Question 27: Customers whose status is 'H' are high-status customers. For each high-status customer, find the number of different products that they purchased. The two attributes appearing in your result should be called customerID and numDiffProducts.

But only include a customer in your result if they purchased at least 4 different products. No duplicates should appear in your result.

Answer 27:

```
SELECT c.customerID, COUNT(DISTINCT sal.productID) AS numDiffProducts
FROM Customers c, Sales sal
WHERE c.status = 'H'
      AND c.customerID = sal.customerID
GROUP BY c.customerID
HAVING COUNT(DISTINCT sal.productID) >= 4;
```

Don't need DISTINCT since c.customerID is the GROUP BY attribute, so there's only one group formed for each c.customerID

Question 28: paidPrice and quantity are attributes in the Sales table. The cost of a sale in Sales equals paidPrice times quantity.

For each store in Stores, give the storeID, storeName, the number of sales for that store and the total of the costs of the Sales for that Store. The last two attributes should appear in your result as numSales and totalCost. No duplicates should appear in your result.

Note that if there are no sales in Sales for a particular store, there your result should include a tuple for that store with storeID, storeName, and with numSales and totalCost both 0. If your result does not provide this, then you'll lose 2 points.

Answer 28: Here are two correct solutions. DISTINCT is not needed in either answer, since no storeID can appear twice in the result of either answer.

```
SELECT st.storeID, st.storeName,  
       COUNT(*) AS numSales,  
       SUM(sal.paidPrice*sal.quantity) AS totalCost  
FROM Stores st, Sales sal  
WHERE st.storeID = sal.storeID  
GROUP BY st.storeID, st.storeName  
  
UNION  
  
SELECT st.storeID, st.storeName,  
       0 AS numSales,  
       0 AS totalCost  
FROM Stores st  
WHERE NOT EXISTS ( SELECT *  
                   FROM Sales sal  
                   WHERE st.storeID = sal.storeID );
```

Or we could use LEFT OUTER JOIN as follows:

```
SELECT st.storeID, st.storeName,  
       COUNT(*) AS numSales,  
       SUM( COALESCE(sal.paidPrice,0) * COALESCE(sal.quantity,0) ) AS totalCost  
FROM Stores st LEFT OUTER JOIN Sales sal  
WHERE st.storeID = sal.storeID  
GROUP BY st.storeID, st.storeName;
```

In both solutions, instead of COUNT(*) it would be okay to use COUNT(st.storeID), or COUNT(X) where X is any attribute of the relations Stores and Sales that cannot be NULL. And since the description of the schema says that no attribute can be NULL, you could use COUNT(X) for any attribute X that's in the relations Stores and Sales. But COUNT(DISTINCT X) would not be correct.

Question 29: This question has two parts; be sure to answer both.

29a): In Products, there are attributes normalPrice and discount. The discounted price of a product equals $\text{normalPrice} - \text{normalPrice} * \text{discount} / 100.00$

Create a view named `ProductsDiscounted`. For each product, this view should provide `productID`, `manufacturer` (`manuf`) and the discounted price of the product. The attributes in your view should appear as `productID`, `manuf` and `discountedPrice`. No duplicates should appear in your result.

Answer 29a):

```
CREATE VIEW ProductsDiscounted AS
SELECT pr.productID, pr.manuf,
       pr.normalPrice - pr.normalPrice * pr.discount / 100 AS discountedPrice
FROM Products pr;
```

DISTINCT is not needed, since productid is the key of Products.

29b): Products from the same manufacturer might have different discounted prices, but there also could be multiple products that have the same discounted price.

Write a query using the ProductsDiscounted view that gives the productID, manufacturer and discountedPrice for each product that has the highest discountedPrice of any product made by that product's manufacturer. If there are multiple products from a manufacturer that have the same highest discountedPrice, then there should be a tuple in your result for each of those products.

Answer 29b): Here are two correct solutions. DISTINCT is not needed in either of them. Some of the incorrect ways of writing queries like this are on Slide 28 of Lecture 5. Aggregates can't be used in the WHERE clause.

[illegible][illegible]

**This is extra blank space, in case you need more paper to answer questions.
Write Question number clearly if you use blank pages.**

You may tear off this page to help you do Part III of the Final.

Some familiar tables appear below, with Primary Keys underlined. **These tables also appear on the last page of the Final, which you can tear off to help you do questions in Part III of the Final.** You don't have to hand in that last page at the end of the Exam.

Products(productID, productName, manuf, normalPrice, discount)

Customers(customerID, custName, address, joinDate, amountOwed, lastPaidDate, status)

Stores(storeID, storeName, region, address, manager)

Days(dayDate, category)

Sales(productID, customerID, storeID, dayDate, paidPrice, quantity)

Payments(customerID, paidDate, amountPaid, cleared)

Assume that no attributes can be NULL, and that there are no UNIQUE constraints.

Assume Referential Integrity as follows:

- Each productID in Sales appears as a productID in Products.
- Each customerID in Sales appears as a customerID in Customers.
- Each storeID in Sales appears as a storeID in Stores.
- Each dayDate in Sales appears as a dayDate in Days.
- Each customerID in Payments appears as a customerID in Customers.

Write legal SQL queries for Questions 25-28. If you want to create and then use views to answer these questions, that's okay, but views are not required unless the question asks for them.

Don't use DISTINCT in your queries unless it's necessary; 1 point will be deducted if you do. Some points will also be deducted for unnecessarily complex queries, including use of GROUP BY when it's not needed.